

Public Railway RAGAS Attempt Report

Generated: 2026-06-22

Outcome

The full RAGAS scoring job was launched against the consolidated 148-case public Railway evaluation output. RAGAS processed all 528 internal metric tasks, but OpenAI judge calls returned insufficient quota errors. As a result, the RAGAS metric columns were all null and no valid aggregate RAGAS score could be computed.

Public Evaluation Context

Signal	Value
Total cases	148
Answerable cases submitted to RAGAS	132
Unanswerable cases retained	16
Citation coverage	1.0000
Source return rate	1.0000
Average sources per case	3.00

RAGAS Result Availability

Metric	Aggregate score	Non-null case scores	Null case scores
faithfulness	n/a	0	132
answer_relevancy	n/a	0	132
context_precision	n/a	0	132
context_recall	n/a	0	132

Error Evidence

Signal	Value
Completed all internal RAGAS tasks	True
Executor exceptions logged	528
RateLimitError occurrences logged	6
insufficient_quota occurrences logged	12

Interpretation

The key issue is not the n8n pipeline or the retrieved sources. The public endpoint had already completed all 148 cases and returned sources for every case. The blocker is judge-side OpenAI quota for RAGAS scoring.

Next Step

Enable billing/quota for the OpenAI judge key or provide a different judge key with available quota, then rerun `score_public_ragas.py`. The script reads the already collected public evaluation output, so it does not need to call the public n8n webhook again.